

Optimized Information Gathering in Robotics and Sensor Networks

Andreas Krause

California Institute of Technology



rsrg@caltech

..where theory and practice collide

Information gathering problems

- Want to **learn something** about the state of the world
 - Estimate water quality in a geographic region, detect outbreaks, ...
- We can choose (partial) **observation**
 - Make measurements, place sensors, choose experimental parameters ...
- ... but they are **expensive / limited**
 - hardware cost, power consumption, grad student time ...



Want to **cost-effectively** get **most useful** information!

Related work

Sensing problems considered in

Experimental design (Lindley '56, Robbins '52...), Spatial statistics

(Cressie '91, ...), Machine Learning (MacKay '92, ...), Robotics (Sim&Roy '05, ...), Sensor Networks (Zhao et al '04, ...), Operations Research (Nemhauser '78, ...)


Existing algorithms typically

- **Heuristics:** No guarantees! Can do arbitrarily badly.
- **Find optimal solutions** (Mixed integer programming, POMDPs):

Very difficult to scale to bigger problems.

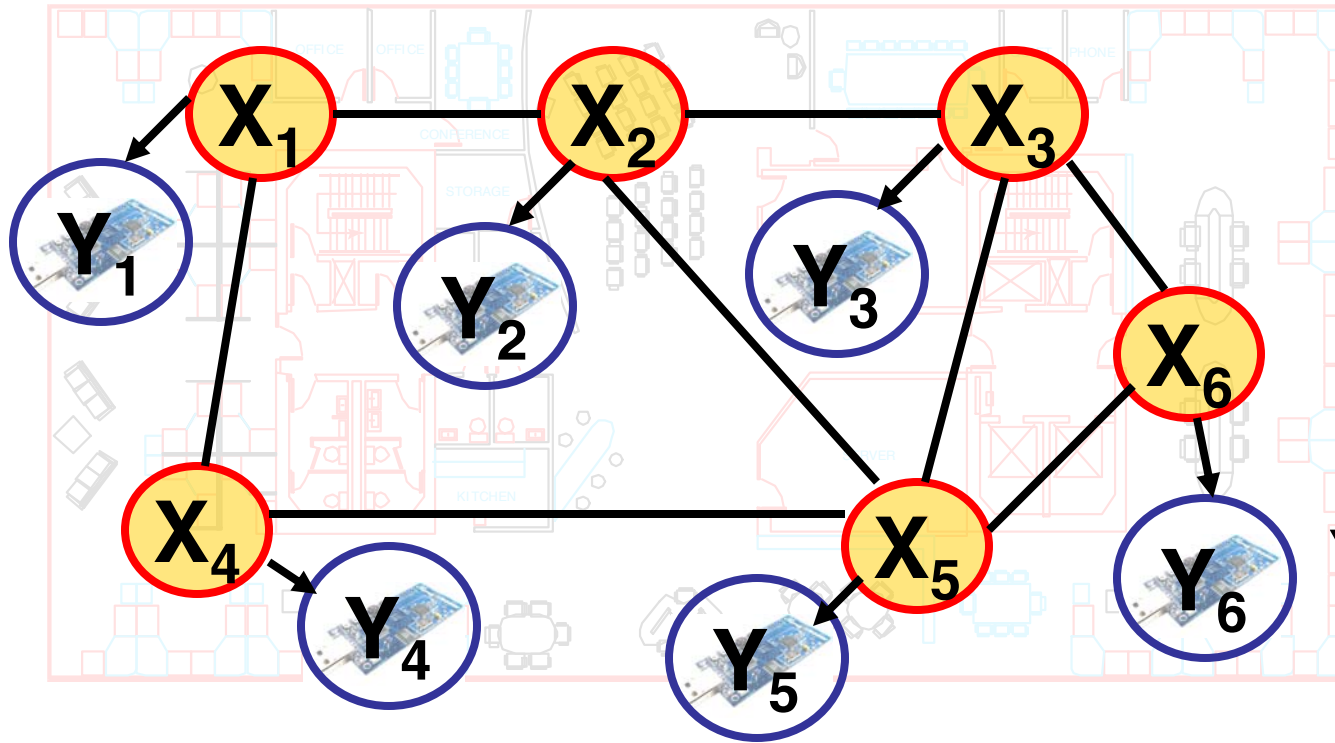
Want algorithms that have theoretical guarantees and scale to large problems!

Running example. Detecting fires



Want to place sensors to detect fires in buildings

Monitoring as Bayesian regression



X_s : temperature at location s

Y_s : sensor value at location s

$$Y_s = X_s + \text{noise}$$


Joint probability distribution

$$P(X_1, \dots, X_n, Y_1, \dots, Y_n) = P(X_1, \dots, X_n) P(Y_1, \dots, Y_n \mid X_1, \dots, X_n)$$

Prior

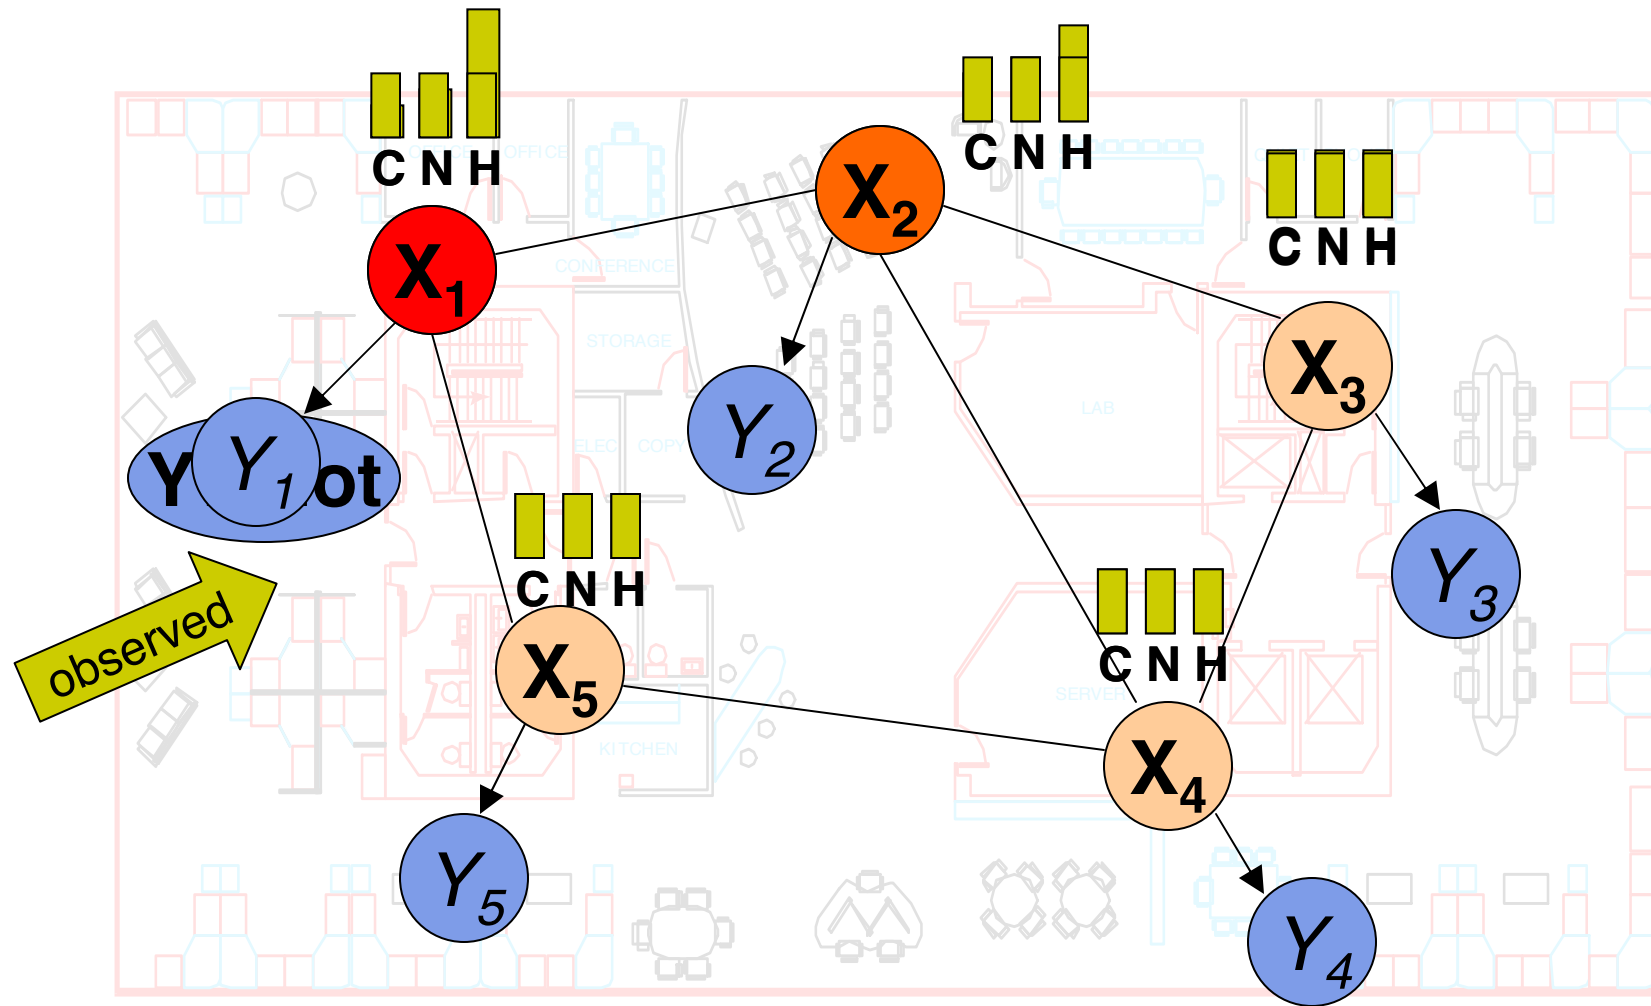
Likelihood

Why is this useful?



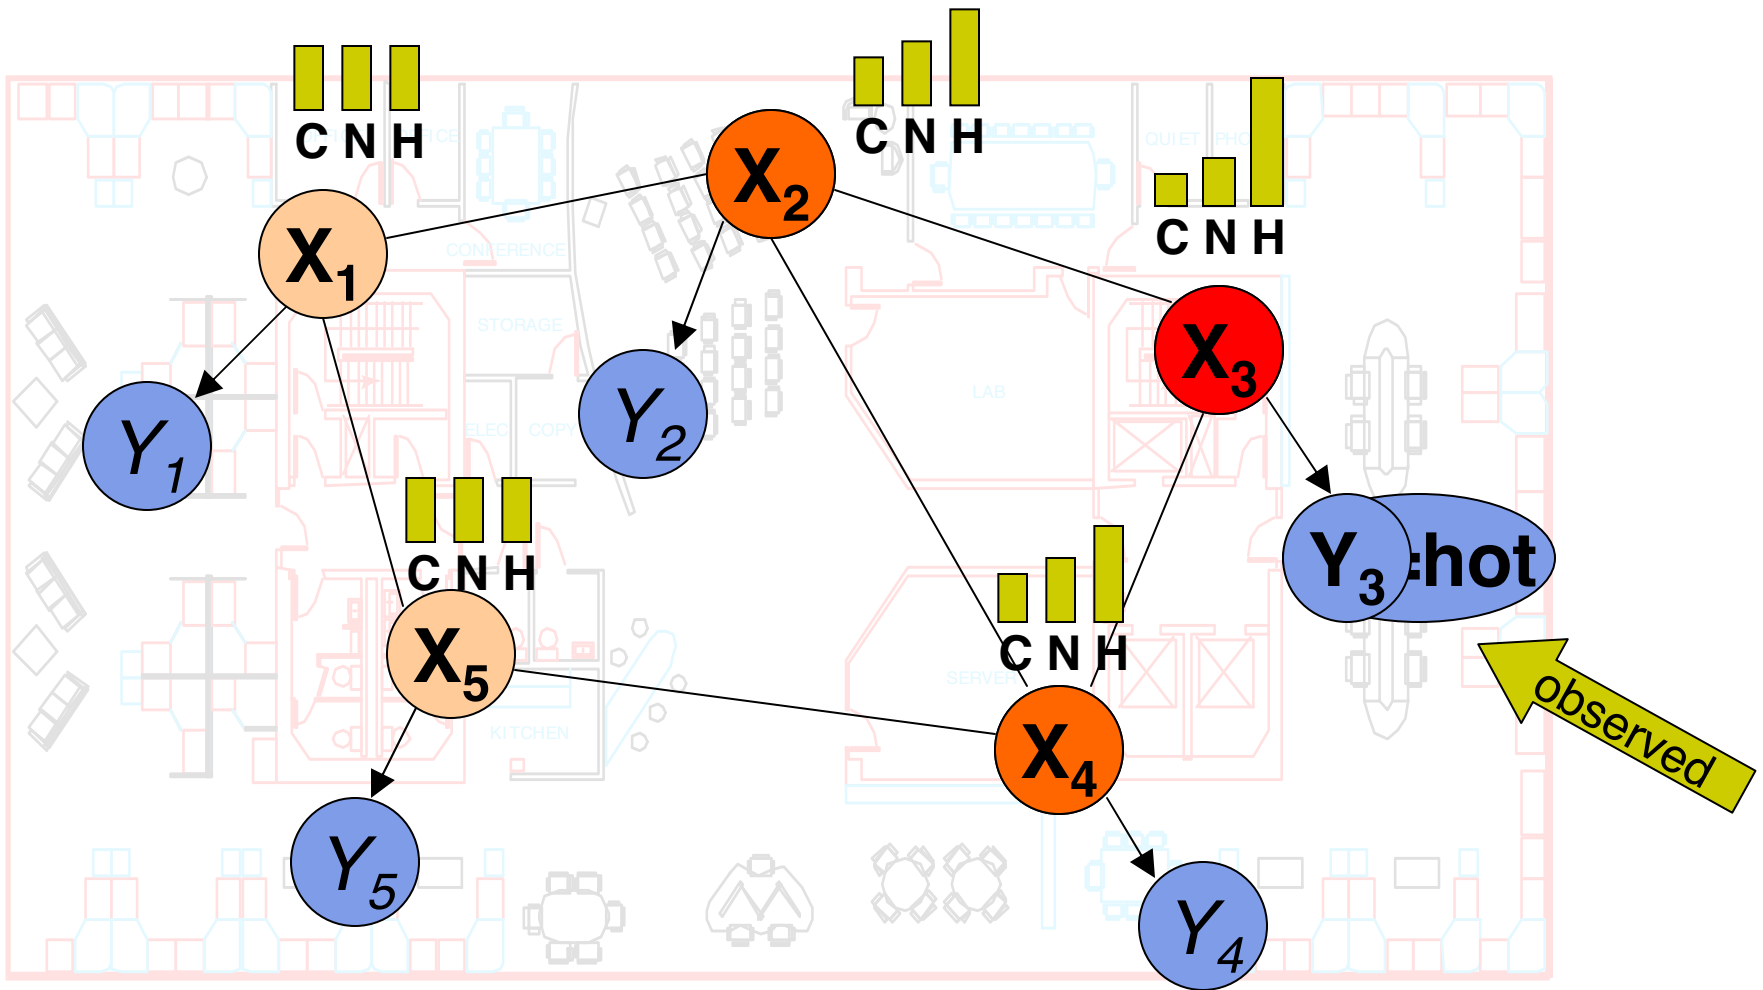
- **Robust reasoning:** Integrate measurements from multiple sensors. E.g.: $P(X_2 \mid y_1, y_2, y_3)$ likely more accurate than $P(X_2 \mid y_2)$
- **Exploiting correlation:** Can predict $P(X_1, X_3 \mid y_2)$
→ Can turn some sensors off to save battery life

Making observations



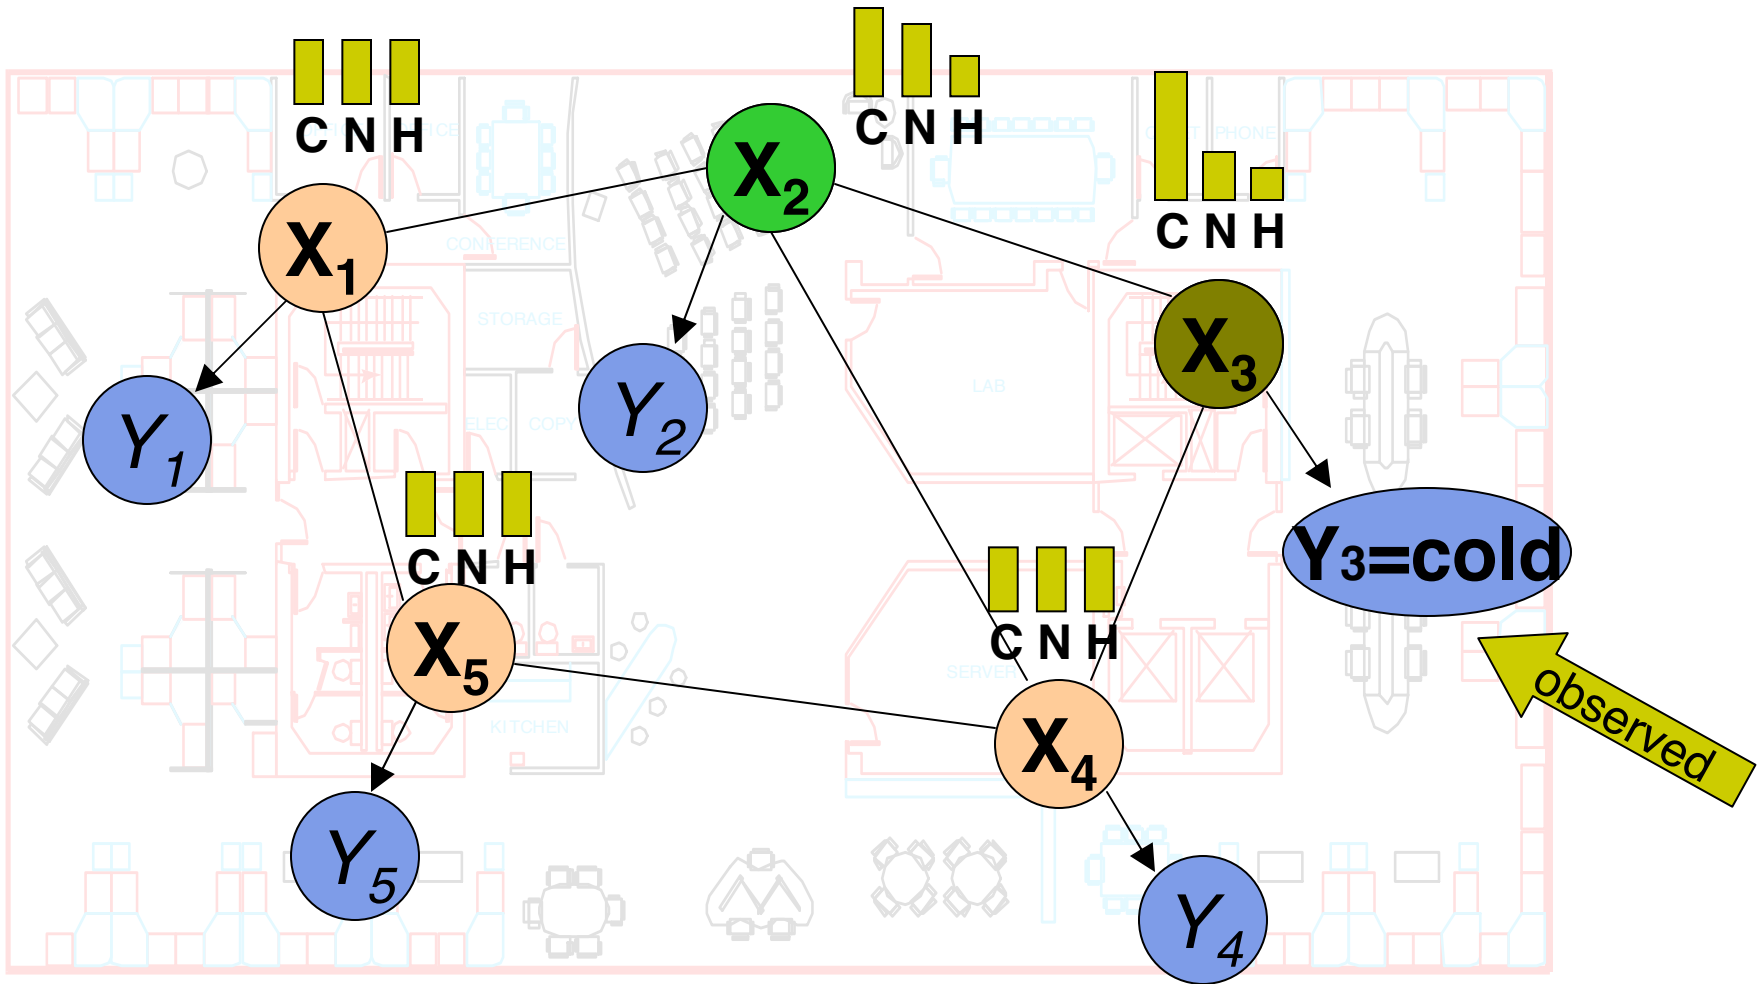
Less uncertain \rightarrow Reward[$P(\mathbf{X} | Y_1 = \text{hot})$] = 0.2

Making observations



$$\text{Reward}[P(\mathbf{X} | Y_3 = \text{hot})] = 0.4$$


A different outcome...



$$\text{Reward}[P(\mathbf{X} | Y_3 = \text{cold})] = 0.1$$

Value of information

Should we raise a fire alert?



Temp. X Actions	<i>Fiery hot</i>	<i>normal/cold</i>
No alarm	-\$\$\$	0
Raise alarm	\$	-\$

Only have belief about temperature $P(X = \text{hot} \mid \text{obs})$

→ choose $a^* = \operatorname{argmax}_a \sum_x P(\mathbf{x}|\text{obs}) U(\mathbf{x}, a)$

Decision theoretic value of information

$\text{Reward}[P(X \mid \text{obs})] = \max_a \sum_x P(\mathbf{x}|\text{obs}) U(\mathbf{x}, a)$

Other example reward functions

Entropy

$$\text{Reward}[P(\mathbf{X})] = -H(\mathbf{X}) = \sum_{\mathbf{x}} P(\mathbf{x}) \log_2 P(\mathbf{x})$$

Expected mean squared prediction error (EMSE)

$$\text{Reward}[P(\mathbf{X})] = -1/n \sum_s \text{Var}(X_s),$$

Many other objectives possible and useful...

Value of information [Lindley '56, Howard '64]

For any set A of sensors, its value of information is

$$F(A) = \sum_{\mathbf{y}_A} \underbrace{P(\mathbf{y}_A)}_{\text{Observations made by sensors } \mathbf{A}} \underbrace{\text{Reward}[P(\mathbf{X} | \mathbf{y}_A)]}_{\text{Reward when observing } Y_A = \mathbf{y}_A}$$

Observations made by sensors \mathbf{A}

Reward when observing $Y_A = \mathbf{y}_A$

Want to find a set $A^* \mu V, |A^*| \leq k$ s.t.

$$A^* = \operatorname{argmax}_{|A| \leq k} F(A)$$

Optimizing Value of Information

- Given: finite set V of locations

- Want: $A^* \subseteq V$ such that
$$A^* = \operatorname{argmax}_{|A| \leq k} F(A)$$

Typically NP-hard!


Greedy algorithm:

Start with $A = \emptyset$;

For $i = 1$ to k



$s^* := \operatorname{argmax}_s F(A \cup \{s\})$

$A := A \cup \{s^*\}$



How well can this simple heuristic do?

Performance of greedy




Temperature data
from sensor network


- Greedy empirically close to optimal. Why?

Key observation: Diminishing returns

Placement A = $\{Y_1, Y_2\}$



Placement B = $\{Y_1, \dots, Y_5\}$



Theorem [Krause and Guestrin, UAI '05]:
Information gain $F(A) = H(X) - H(X \mid Y_A)$ is
submodular!

New sensor Y'

Submodularity: 

+ $\bullet Y'$  Large improvement

+ $\bullet Y'$  Small improvement

For $A \mu B$, $F(A \cup \{Y'\}) - F(A) \geq F(B \cup \{Y'\}) - F(B)$

One reason submodularity is useful

Theorem [Nemhauser et al '78]

Greedy algorithm gives constant factor approximation

$$F(A_{\text{greedy}}) \geq (1 - 1/e) F(A_{\text{opt}})$$

~63%

- Greedy algorithm gives near-optimal solution!
- For information gain: Guarantees best possible unless $P = NP$!
[Krause & Guestrin '05]

Challenges for environmental monitoring



Use robots to monitor environment




Not just select best k locations A for given $F(A)$. Need to

... be **robust** against uncertainty in the function F


... take into account **cost** of traveling between locations

... cope with environments that **change** over time

Why do we need robustness?



→ Sensor failures




→ Adversarial environments

$$\text{Unified view: } \mathcal{A}^* = \underset{|\mathcal{A}| \leq k}{\operatorname{argmax}} \min_i F_i(\mathcal{A})$$

How does the greedy algorithm do?

$V = \{s_1, s_2, s_3\}$ Buy $k=2$ sensors $F_i = \text{intrusion at } s_i$



Optimal
solution

Optimal score: 1

Greedy picks
 s_3 first

Then, can
choose only
 s_1 or s_2

Greedy score: 2

→ Greedy does arbitrarily badly 😞

Can we do better?